

Standard data-quality checks flagged 4 bad rows. A structural check flagged 1,931.

52% of the dataset was duplicated — and every duplicated row passed validation individually.

3,732 → 1,801 ROWS → REAL PRODUCTS 1,931 redundant rows removed	14 → 9 CATEGORIES 5 were duplicate blocks	+112% INVENTORY VALUE ERROR ₹22.4L reported vs ₹10.6L actual	28.6% WORST AVAILABILITY Biscuits, vs 12.0% average
--	--	---	--

WHAT EVERY STANDARD CHECK FOUND

Blank name or category	0
MRP = 0	1
Selling price = 0	1
Selling price > MRP	0
Weight = 0	2
Out-of-stock flag vs stock on hand	0
Stated discount % vs price maths	0
Total rows failing	4

By every conventional measure, a clean dataset.

WHAT THE DUPLICATION DID TO THE HEADLINE NUMBERS

METRIC	REPORTED	ACTUAL	ERROR
Total inventory value	₹2,243,081	₹1,058,454	+112%
Product records	3,732	1,801	+107%
SKUs >₹300 out of stock	8	4	+100%
>₹500 MRP, <10% discount	82	39	+110%
Top-5 avg-discount table	3 of 5 entries were phantom categories		—

No row-level rule detects this. The problem is the shape of the dataset, not the values inside it.

WHAT COMPARING BLOCK SHAPES FOUND

Cooking Essentials	= Munchies	514/514
Personal Care	= Paan Corner	344/344
Packaged Food	= Ice Cream & Desserts	388/388
Packaged Food	= Chocolates & Candies	388/388
Dairy, Bread & Batter	= Beverages	129/129

Byte-identical, position for position. Semantics confirm the first block holds the correct label: *paneer* appears under Dairy, then again under Beverages.

CORRECTED VIEW: AVAILABILITY BY CATEGORY

Biscuits		28.6%
Dairy & Bread		23.9%
Meats, Fish, Eggs		21.9%
Cooking Essentials		12.5%
Home & Cleaning		9.8%
Packaged Food		8.9%
Fruits & Veg		6.5%
Health & Hygiene		6.5%
Personal Care		6.1%

Catalogue average 12.0%. Biscuits runs at ~5× Personal Care — a replenishment problem in a fast-moving, low-value category, not a pricing one.

WHERE THE VALUE SITS

Cooking Essentials		₹337k
Personal Care		₹270k
Packaged Food		₹192k
Home & Cleaning		₹123k
Other 5 categories		₹136k

Two categories hold 57% of ₹1,058,454 of stock. Inventory value at selling price — not revenue.

HIGH-VALUE SKUS UNAVAILABLE

PRODUCT	MRP
Patanjali Cow's Ghee	₹565
MamyPoko Pants Standard, XL	₹399
Aashirvaad Atta w/ Multigrains	₹315
Everest Kashmiri Lal Chilli	₹310

Four, not eight. Eight looks like a pattern; four is a shortlist someone can action this week.

DISCOUNT POSITION

Catalogue average discount	7.62%
Fruits & Vegetables (deepest)	15.46%
Personal Care (shallowest)	6.24%
Margin given away, stock on hand	₹115,512
as % of MRP value	9.8%

F&V discounts ~2× the average — consistent with perishability, not promotion.

Quarantine, don't delete

DELETE FROM zepto WHERE mrp = 0 destroys the evidence before you've counted it. Failed rows go to zepto_rejects with a reason, so "how many failed and why" stays a query.

No CHECK that can't fire

A CHECK (dsp <= mrp) guarantees the matching diagnostic returns zero rows — dead code that reads as a passing test. If the raw data does violate it, the whole INSERT aborts. Validation belongs in the quarantine step.

Inventory value ≠ revenue

Selling price × quantity on hand is what's in the warehouse. Revenue needs units sold, which this dataset has no column for. Also: discount_percent is FLOOR, not ROUND — every discount reads up to 1pp low.